

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

$$\mathbf{P}\mathbf{A} = \mathbf{L}\mathbf{U}$$

$$\mathbf{A}\mathbf{v} = \lambda\mathbf{v}, \quad \mathbf{A}^\top\mathbf{A}\mathbf{v} = \sigma^2\mathbf{v}$$

Juliaと線形代数の基礎：ソルバー選定とコード性能最適化

王 安清（ワン アンチン）

東北大学 環境科学研究所 福島研

研究分野：プロセスシステム工学

（設計・プロセス最適化）

```
using LinearAlgebra, SparseArrays
x = A \ b
F = cholesky(Symmetric(A))
ldiv!(x, F, b)
```

$$\mathbf{A}\mathbf{x} = \mathbf{b}$$

$$\begin{bmatrix} \mathbf{B}_1 & \mathbf{C}_1 \\ \mathbf{A}_2 & \mathbf{B}_2 & \mathbf{C}_2 \\ \ddots & \ddots & \ddots \\ & & \mathbf{A}_n & \mathbf{B}_n \end{bmatrix}$$

JuliaLang Japan 2025

2025年12月13日

東京科学大学 大岡山キャンパス

Overview

- Why learn linear solvers?
- How to choose a solver for $Ax = b$
 - A is symmetric/Hermitian
 - A nonsymmetric/non-Hermitian
- Example 1: Thomas algorithm
- Example 2: Block thomas algorithm

Why learn linear solvers?

- Solving linear systems is a fundamental task in engineering: least squares, linear regression, partial differential equation (PDE) models, optimization (KKT systems), nonlinear solvers, etc.
- In practice, we must choose algorithms that exploit the matrix structure and strike a good balance between numerical accuracy and performance.

Why learn linear solvers?

- Accuracy.

$$\mathbf{A} = \begin{bmatrix} 10^{-20} & 1 \\ 1 & 1 \end{bmatrix}, \mathbf{b} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$$

$$\mathbf{A} = \begin{bmatrix} 1 & 10^{20} \\ 0 & 1 \end{bmatrix}, \mathbf{b} = \begin{bmatrix} 10^{20} \\ 1 \end{bmatrix}$$

→ Gaussian elimination

$$\mathbf{A} = \begin{bmatrix} 1 & 10^{20} \\ 1 & 1 \end{bmatrix}, \mathbf{b} = \begin{bmatrix} 10^{20} \\ 2 \end{bmatrix}$$

$$\mathbf{A} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \mathbf{b} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

$\mathbf{x} \approx \begin{bmatrix} 0 \\ 1 \end{bmatrix}$ Gaussian elimination with partial pivoting (GEPP)

$$\mathbf{A} = \begin{bmatrix} 1 & 10^{20} \\ 0 & -10^{20} \end{bmatrix}, \mathbf{b} = \begin{bmatrix} 10^{20} \\ -10^{20} \end{bmatrix}$$

$$\mathbf{x}^* \approx \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

Why learn linear solvers?

- Performance.

$$\begin{bmatrix} 2 & -1 & 0 & \cdots & 0 \\ -1 & 2 & -1 & \ddots & \vdots \\ 0 & -1 & \ddots & \ddots & 0 \\ \vdots & \ddots & \ddots & 2 & -1 \\ 0 & \cdots & 0 & -1 & 2 \end{bmatrix}$$

```
using LinearAlgebra, BenchmarkTools
n = 10_000
d = fill(2.0, n)
dl = fill(-1.0, n - 1)
du = fill(-1.0, n - 1)
Atri = Tridiagonal(dl, d, du)
b = ones(n)

# Dense LU (blocked GEPP) (O(n3) time, O(n2) memory)
Adense = Matrix(Atri)
@btime $Adense \ $b;
```

969.285 ms (9 allocations: 763.09 MiB)

```
# Tridiagonal solve (O(n) time, O(n) memory)
@btime $Atri \ $b;
```

129.100 μs (20 allocations: 469.31 KiB)

Later, we'll show how to write an algorithm that further reduces the runtime to **44.600 μs** with **zero allocations (0 bytes)**.

Choosing a solver for $Ax = b$

- When solving $Ax = b$, the first thing to check is:
Is A symmetric or Hermitian?

- Symmetric matrix.**

A real square matrix that equals its transpose:

$$A = \begin{bmatrix} 4 & 1 & 0 \\ 1 & 3 & -1 \\ 0 & -1 & 2 \end{bmatrix}, \quad A^T = A$$

- Hermitian matrix.**

A complex square matrix that equals its conjugate transpose:

$$B = \begin{bmatrix} 2 & 1+i \\ 1-i & 3 \end{bmatrix}, \quad B^H = \overline{B}^T = B$$

For symmetric/Hermitian matrix A

- Consider whether A is Strictly Positive Definite (SPD):

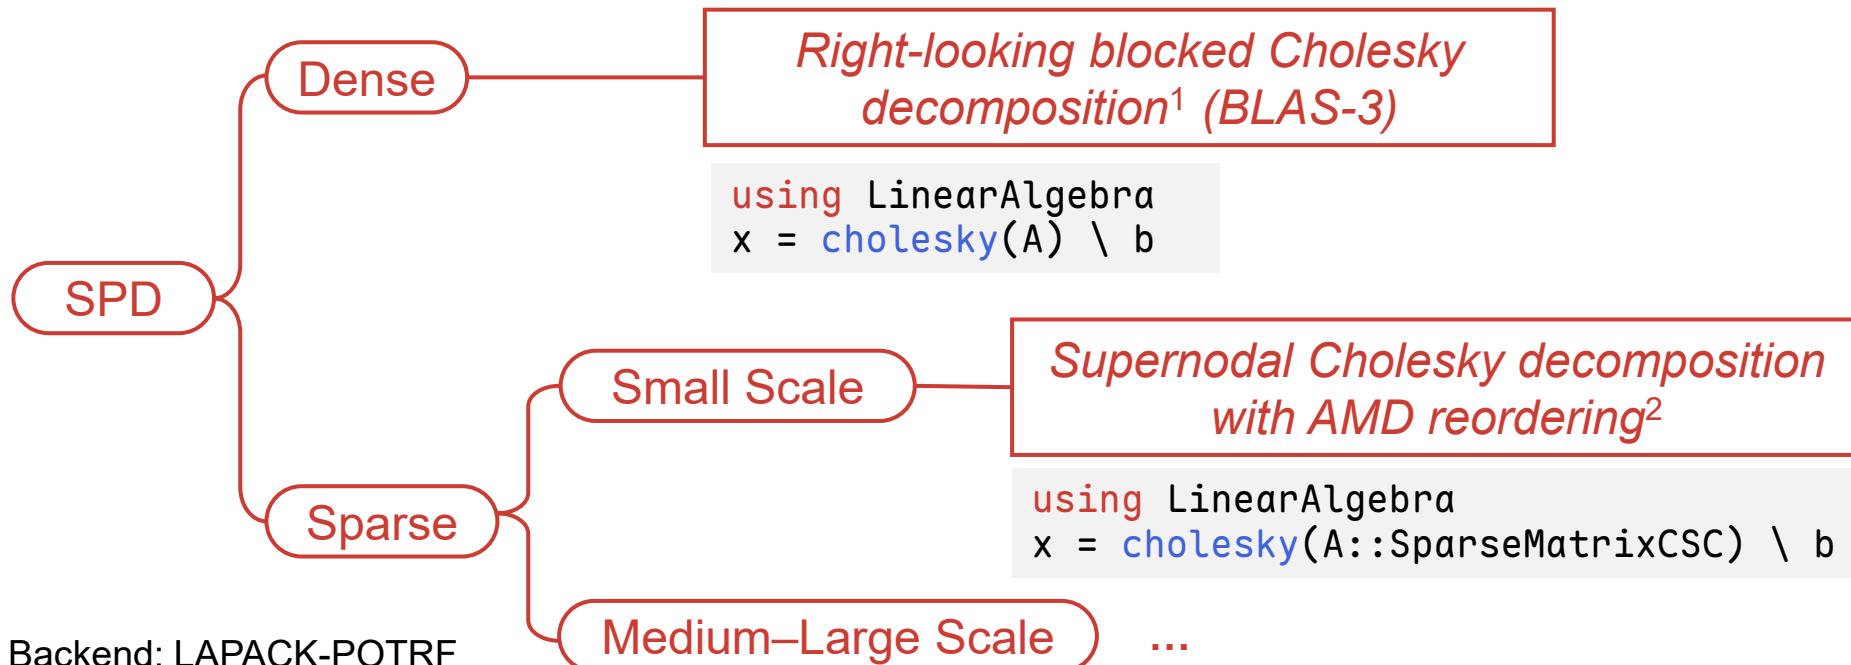
$$A = A^T, \quad x^T A x > 0 \quad \forall x \neq 0$$

$$\begin{bmatrix} 2 & -1 \\ -1 & 2 \end{bmatrix}$$

$$\begin{bmatrix} 2 & -1 & 0 \\ -1 & 2 & -1 \\ 0 & -1 & 2 \end{bmatrix}$$

- Equivalently:

$$\lambda_i(A) > 0 \text{ for all } i;$$



¹ Backend: LAPACK-POTRF

² Backend: SuiteSparse-CHOLMOD

Large scale sparse SPD

- Conjugate Gradient (CG)

```
using IterativeSolvers
x, [history] = IterativeSolvers.cg(A, b; kwargs...)

using Krylov
x, stats = Krylov.cg(A, b::AbstractVector{FC};)

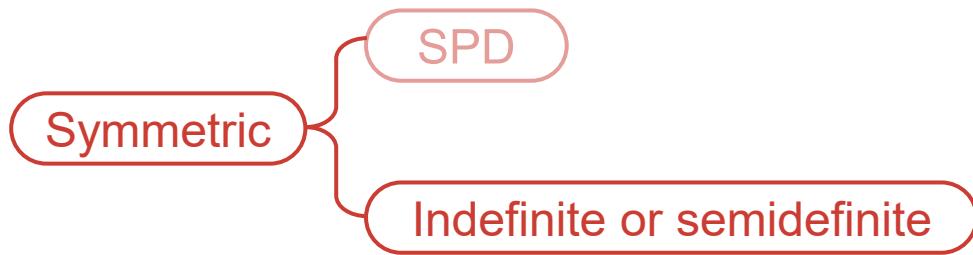
using LinearSolve # SciML unified ecosystem
sol = solve(LinearProblem(A, b; u0=x0), KrylovJL_CG())
sol = solve(LinearProblem(A, b; u0=x0), IterativeSolversJL_CG())
```

- Optional: build a preconditioner from A .

- Related packages: Preconditioners.jl, IncompleteCholesky.jl, LinearAlgebra
- Example (Jacobi/diagonal preconditioner)

```
using Krylov, LinearAlgebra
M = Diagonal(diag(A))
x, stats = Krylov.cg(A, b; M)
```

Indefinite or semidefinite



- Small Scale
 - Bunch-Kaufman pivoted LDL^T factorization¹

```
using LinearAlgebra
x = bunchkaufman(A) \ b
```

- Large Scale
 - MINimum RESidual method (MINRES)

```
using IterativeSolvers
x, [history] = IterativeSolvers.minres(A, b; kwargs...)
using Krylov
x, stats = Krylov.minres(A, b::AbstractVector)
```

¹ Backend: SuiteSparse-CHOLMOD

For nonsymmetric/non-Hermitian matrix A

- Small scale, dense
 - Blocked gaussian elimination with partial pivoting (GEPP) (LU factorization)¹

```
using LinearAlgebra  
x = lu(A::AbstractMatrix) \ b
```

- Small scale, sparse
 - Unsymmetric multifrontal LU factorization with partial pivoting²

```
using LinearAlgebra, SparseArrays  
x = lu(A::AbstractSparseMatrixCSC) \ b
```

¹ Backend: LAPACK-GETRF

² Backend: SuiteSparse-UMFPACK

For nonsymmetric/non-Hermitian matrix A

- Large scale
 - Unsymmetric multifrontal LU factorization with partial pivoting
 - Generalized Minimal Residual (GMRES)

```
using IterativeSolvers
x, [history] = IterativeSolvers.gmres(A, b; kwargs...)

using Krylov
x, stats = Krylov.gmres(A, b)
```

- Bi-Conjugate Gradient Stabilized (BiCGSTAB)

```
using Krylov
x, stats = Krylov.bicgstab(A, b)
```

For nonsymmetric/non-Hermitian matrix A

- Large scale
 - Quasi-Minimal Residual (QMR) method

```
using IterativeSolvers
x, [history] = IterativeSolvers.qmr(A, b; kwargs...)

using Krylov
x, stats = Krylov.qmr(A, b)
```

- ...
- Preconditioner
 - IncompleteLU.jl

Special structures

- SPD banded → Banded Cholesky $A = LL^\top$

$$A = \begin{bmatrix} 4 & -1 & 0 & 0 \\ -1 & 4 & -1 & 0 \\ 0 & -1 & 4 & -1 \\ 0 & 0 & -1 & 4 \end{bmatrix}$$

- Symmetric indefinite banded → Banded pivoted LDL^\top

$$A = \begin{bmatrix} 2 & -1 & 0 & 0 \\ -1 & 0 & -1 & 0 \\ 0 & -1 & 2 & -1 \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

- Nonsymmetric banded → Banded LU with partial pivoting

$$A = \begin{bmatrix} 4 & -1 & 2 & 0 & 0 & 0 \\ -2 & 5 & -1 & 1 & 0 & 0 \\ 0 & -3 & 6 & -1 & 2 & 0 \\ 0 & 0 & -2 & 5 & -1 & 1 \\ 0 & 0 & 0 & -3 & 5 & -1 \\ 0 & 0 & 0 & 0 & -2 & 4 \end{bmatrix}$$

Special structures

- Tridiagonal (general) → Thomas algorithm

$$\mathbf{A} = \begin{bmatrix} \beta_1 & \gamma_1 & & & \\ \alpha_2 & \beta_2 & \gamma_2 & & \\ & \ddots & \ddots & \ddots & \\ & & \alpha_n & \beta_n & \end{bmatrix}$$

- Block tridiagonal → Block Thomas / block LU

$$\mathbf{A} = \begin{bmatrix} \mathbf{B}_1 & \mathbf{C}_1 & & & \\ \mathbf{A}_2 & \mathbf{B}_2 & \mathbf{C}_2 & & \\ & \ddots & \ddots & \ddots & \\ & & \mathbf{A}_m & \mathbf{B}_m & \end{bmatrix}$$

- Cyclic tridiagonal → Sherman–Morrison (SMW) or Cyclic Reduction (CR)

$$\mathbf{A} = \begin{bmatrix} \beta_1 & \gamma_1 & 0 & \cdots & \alpha_1 \\ \alpha_2 & \beta_2 & \gamma_2 & \ddots & 0 \\ 0 & \ddots & \ddots & \ddots & 0 \\ \vdots & \ddots & \alpha_{n-1} & \beta_{n-1} & \gamma_{n-1} \\ \gamma_n & 0 & \cdots & \alpha_n & \beta_n \end{bmatrix}$$

Example 1: Thomas algorithm

- We solve the tridiagonal system with $a_1 = 0, c_n = 0$.

$$\left[\begin{array}{cccccc|c} b_1 & c_1 & 0 & 0 & 0 & 0 & d_1 \\ a_2 & b_2 & c_2 & 0 & 0 & 0 & d_2 \\ 0 & a_3 & b_3 & c_3 & 0 & 0 & d_3 \\ 0 & 0 & a_4 & b_4 & c_4 & 0 & d_4 \\ 0 & 0 & 0 & a_5 & b_5 & c_5 & d_5 \\ 0 & 0 & 0 & 0 & a_6 & b_6 & d_6 \end{array} \right]$$

- Step 1. Forward sweep**
- Eliminate subdiagonal entries.

$$\left[\begin{array}{cccccc|c} 1 & c'_1 & 0 & 0 & 0 & 0 & d'_1 \\ 0 & \tilde{b}_2 & c_2 & 0 & 0 & 0 & d'_2 \\ 0 & a_3 & b_3 & c_3 & 0 & 0 & d_3 \\ 0 & 0 & a_4 & b_4 & c_4 & 0 & d_4 \\ 0 & 0 & 0 & a_5 & b_5 & c_5 & d_5 \\ 0 & 0 & 0 & 0 & a_6 & b_6 & d_6 \end{array} \right]$$

- Step 2. Back substitution**
- Now the system is upper triangular in terms of c'_i, d'_i .
- $x_i = d'_i - c'_i x_{i+1}$.

$$\left[\begin{array}{cccccc|c} 1 & c'_1 & 0 & 0 & 0 & 0 & d'_1 \\ 0 & 1 & c'_2 & 0 & 0 & 0 & d'_2 \\ 0 & 0 & 1 & c'_3 & 0 & 0 & d'_3 \\ 0 & 0 & 0 & 1 & c'_4 & 0 & d'_4 \\ 0 & 0 & 0 & 0 & 1 & c'_5 & d'_5 \\ 0 & 0 & 0 & 0 & 0 & 1 & d'_6 \end{array} \right]$$

Example 1: Thomas algorithm

```
function thomas_tridiagonal(
    L::AbstractVector, # subdiagonal (a)
    D::AbstractVector, # diagonal      (b)
    U::AbstractVector, # superdiag    (c)
    b::AbstractVector  # right-hand side
)
    n      = length(b)
    Upos = Vector{Float64}(undef, n - 1) # modified superdiagonal (c')
    x     = Vector{Float64}(undef, n)

    # Forward sweep
    Upos[1] = U[1] / D[1]
    x[1]   = b[1] / D[1]
    for i in 2:n-1
        di      = D[i] - L[i] * Upos[i-1]
        Upos[i] = U[i] / di                      # c'_i
        x[i]   = (b[i] - L[i] * x[i-1]) / di    # d'_i
    end
    di      = D[n] - L[n] * Upos[n-1]
    x[n]   = (b[n] - L[n] * x[n-1]) / di

    # Back substitution
    for i in n-1:-1:1
        x[i] -= Upos[i] * x[i+1]
    end

    return x
end
```

61.600 µs (6 allocations: 156.39 KiB)

Example 1: Thomas algorithm

```
function thomas_tridiagonal(  
    L::AbstractVector, # subdiagonal (a)  
    D::AbstractVector, # diagonal      (b)  
    U::AbstractVector, # superdiag     (c)  
    b::AbstractVector  # right-hand side  
)  
    n = length(b)  
    Upos = Vector{Float64}(undef, n - 1) # modified superdiagonal (c')  
    x = Vector{Float64}(undef, n)  
    @inbounds begin  
        # Forward sweep  
        Upos[1] = U[1] / D[1]  
        x[1] = b[1] / D[1]  
        for i in 2:n-1  
            di = D[i] - L[i] * Upos[i-1]  
            Upos[i] = U[i] / di          # c'_i  
            x[i] = (b[i] - L[i] * x[i-1]) / di  # d'_i  
        end  
        di = D[n] - L[n] * Upos[n-1]  
        x[n] = (b[n] - L[n] * x[n-1]) / di  
  
        # Back substitution  
        for i in n-1:-1:1  
            x[i] -= Upos[i] * x[i+1]  
        end  
    end  
    return x  
end
```

61.600 µs (6 allocations: 156.39 KiB)
46.900 µs (5 allocations: 156.36 KiB)

Example 1: Thomas algorithm

```
function thomas_tridiagonal!(x::AbstractVector,
    L::AbstractVector, # subdiagonal (a)
    D::AbstractVector, # diagonal (b)
    U::AbstractVector, # superdiag (c)
    b::AbstractVector # right-hand side
)
    n      = length(b)
    Upos = Vector{Float64}(undef, n - 1) # modified superdiagonal (c')
    @inbounds begin
        # Forward sweep
        Upos[1] = U[1] / D[1]
        x[1]   = b[1] / D[1]
        for i in 2:n-1
            di      = D[i] - L[i] * Upos[i-1]
            Upos[i] = U[i] / di          # c'_i
            x[i]   = (b[i] - L[i] * x[i-1]) / di  # d'_i
        end
        di      = D[n] - L[n] * Upos[n-1]
        x[n]   = (b[n] - L[n] * x[n-1]) / di
        # Back substitution
        for i in n-1:-1:1
            x[i] -= Upos[i] * x[i+1]
        end
    end
    return x
end
```



Need buffer!

Global scratch buffer

```
const Upos = Vector{Float64}(undef, n - 1)

function get_buffers(n::Int)
    length(Upos) != n && resize!(Upos, n)
    return Upos
end

function thomas_tridiagonal!(...)
    ...
end
```

- Problem: **Not thread safe**

- Two threads can resize at the same time
- The buffer can be reallocated while another thread is using it
- Leads to data races

Task local storage for scratch buffers

```
using TaskLocalValues

const TLS_buf = TaskLocalValue{Vector{Float64}}(() ->
Float64[])

function get_buffers(n::Int)
    buf = TLS_buf[]
    length(buf) != n && resize!(buf, n)
    return buf
end
```

Example 1: Thomas algorithm

```
function thomas_tridiagonal!(x::AbstractVector,
    L::AbstractVector, # subdiagonal (a)
    D::AbstractVector, # diagonal      (b)
    U::AbstractVector, # superdiag    (c)
    b::AbstractVector  # right-hand side
)
    n      = length(b)
    Upos = get_buffers(n - 1) # modified superdiagonal (c')
    @inbounds begin
        # Forward sweep
        Upos[1] = U[1] / D[1]
        x[1]   = b[1] / D[1]
        for i in 2:n-1
            di      = D[i] - L[i] * Upos[i-1]
            Upos[i] = U[i] / di          # c'_i
            x[i]   = (b[i] - L[i] * x[i-1]) / di  # d'_i
        end
        di      = D[n] - L[n] * Upos[n-1]
        x[n]   = (b[n] - L[n] * x[n-1]) / di
        # Back substitution
        for i in n-1:-1:1
            x[i] -= Upos[i] * x[i+1]
        end
    end
    return x
end
```

61.600 µs (6 allocations: 156.39 KiB)
46.900 µs (5 allocations: 156.36 KiB)
44.600 µs (0 allocations: 0 bytes)

Example 2: Block thomas algorithm

$$A = \begin{bmatrix} D_1 & U_1 & & & \\ L_2 & D_2 & U_2 & & \\ & \ddots & \ddots & \ddots & \\ & & L_{n-1} & D_{n-1} & U_{n-1} \\ & & & L_n & D_n \end{bmatrix} \quad x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad b = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix}$$

$$\tilde{U}_1 = D_1^{-1} U_1,$$

$$\tilde{b}_1 = D_1^{-1} b_1,$$

$$S_i = D_i - L_i \tilde{U}_{i-1},$$

$$\tilde{U}_i = S_i^{-1} U_i,$$

$$\tilde{b}_i = S_i^{-1} (b_i - L_i \tilde{b}_{i-1}), \quad i = 2, \dots, n-1,$$

$$S_n = D_n - L_n \tilde{U}_{n-1},$$

$$\tilde{b}_n = S_n^{-1} (b_n - L_n \tilde{b}_{n-1}).$$

Example 2: Block thomas algorithm

```
function block_thomas_tridiagonal!(x, L, D, U, b)
    n = size(D[1], 1)
    m = length(D)
    lu_buf = Vector{LU{Float64},Matrix{Float64},Vector{Int64}}(undef, m - 1)

    # Forward elimination
    @inbounds for i = 1:m-1
        F = lu_buf[i] = lu!(D[i])
        rdiv!(L[i], F)
        # D[i+1] <- -1*L[i]*U[i] + 1*D[i+1]
        # BLAS GEMM: C = α A B + β C
        mul!(D[i+1], L[i], U[i], -1, 1)
        b_i = @view b[(i-1)*n+1:i*n]
        b_{i+1} = @view b[i*n+1:(i+1)*n]
        # b_{i+1} <- b_{i+1} - L[i] * b_i
        # BLAS GEMV: y = α A x + β y
        mul!(b_{i+1}, L[i], b_i, -1, 1)
    end
end
```

$$\begin{aligned}\tilde{U}_1 &\leftarrow D_1^{-1} U_1, \\ \tilde{b}_1 &\leftarrow D_1^{-1} b_1, \\ D_{i+1} &\leftarrow D_{i+1} - L_i \tilde{U}_i, \\ \tilde{U}_{i+1} &\leftarrow D_{i+1}^{-1} U_{i+1}, \\ \tilde{b}_{i+1} &\leftarrow D_{i+1}^{-1} (b_{i+1} - L_i \tilde{b}_i), \\ &\quad \text{for } i = 1, \dots, n-1.\end{aligned}$$

Example 2: Block thomas algorithm

```
# Backward substitution ( $D[i]*x_i = b_i - U[i]*x_{i+1}$ )
Fm = lu!(D[m])
xm = @view x[(m-1)*n+1:m*n]
bm = @view b[(m-1)*n+1:m*n]
ldiv!(xm, Fm, bm)
@inbounds for i = m-1:-1:1
    xi = @view x[(i-1)*n+1:i*n]
    rhs = @view b[(i-1)*n+1:i*n]
    xi_1 = @view x[i*n+1:(i+1)*n]
    ldiv!(xi, lu_buf[i], mul!(rhs, U[i], xi_1, -1, 1))
end

return x
end
```

Conclusions

- **Solving linear systems**

- Check matrix structure first (tridiagonal, block, banded, sparse, ...)
- Choose an algorithm that matches the structure
- Then pick the Julia library / implementation
- Add a preconditioner when necessary

- **Writing your own solver**

- A custom implementation can sometimes outperform general-purpose libraries.
- Prefer **in-place, mutating operations** in hot loops.
- Use **task-local scratch buffers** to eliminate allocations while remaining thread-safe.
- Whenever possible, **rewrite the mathematics into BLAS-friendly kernels**, rather than copying formulas line by line into code.

ご清聴ありがとうございました
Thank you for your attention

王 安清
GitHub: [@abcdvvvv](https://github.com/abcdvvvv)

Scan to download the slides (PDF) →

